




torchsetlin

A GPU-native, PyTorch-idiomatic framework
for Tsetlin machines

Vajira Thambawita

Simula Metropolitan Center for Digital Engineering, Oslo, Norway

vajira@simula.no

Abstract. The Tsetlin machine is a propositional-logic learner: it builds classifiers out of conjunctive clauses over Boolean literals, updated not by gradient descent but by reward/penalty feedback to teams of Tsetlin automata. Its models are legible by construction and its inference is bit-level cheap, which is why it has become attractive for clinical decision support, edge inference and hardware accelerators. Its software, however, has grown up outside the deep-learning ecosystem: the reference implementations are C and CUDA kernels behind scikit-learn-shaped wrappers, each model family a separate package, with no shared tensor abstraction, no common training loop and no route into PyTorch tooling.

`torchsetlin` closes that gap. It expresses the whole Tsetlin machine family — vanilla, weighted, coalesced, regression and convolutional — as ordinary `torch.nn.Module` objects whose learnable state is a tensor of automaton counters. A model therefore moves across devices with `.to(device)`, serialises with `state_dict()`, and learns from a mini-batch with a single `model.update(x, y)` call that plays the role of `loss.backward(); optimizer.step()`. The central technical contribution is a *batched* formulation of Type I/Type II feedback in which per-example stochastic automaton transitions are replaced by event counts accumulated with two matrix products and realised as binomial draws — one commit per batch instead of one per example. On an RTX 3090 this reaches 2.2×10^5 training examples/s and up to $273\times$ the throughput of the same code on a 16-core CPU, while matching the sequential algorithm's accuracy at moderate batch sizes. Around the models the package ships Booleanization encoders, a Keras-style trainer with callbacks, metrics, closed-form global and local interpretability, and visualisation — everything a Tsetlin machine workflow needs, in one PyTorch-shaped package.

Software: `torchsetlin 0.1.3` | MIT licence | `pip install torchsetlin`

Code: <https://github.com/vlbthambawita/torchsetlin> | **Docs:** <https://vlbthambawita.github.io/torchsetlin/>

Measurements in Section 8 were produced by the benchmark driver shipped with the repository (`benchmarks/run_benchmarks.sh`) and can be reproduced on other hardware with one command.

Contents

1	Executive summary	1
2	Motivation	2
2.1	Why a logic learner at all	2
2.2	Three obstacles in practice	3
3	Existing solutions and what they leave open	4
3.1	The landscape	4
3.2	The gap	4
4	The proposed library	5
4.1	Design principles	5
4.2	What the library provides	6
5	System architecture	7
5.1	Four layers	7
5.2	The learning state	8
5.3	How an update flows	9
5.4	Extension points	10
6	Technical design	10
6.1	Representation	10
6.2	Inference is two matrix products	10
6.3	The feedback policy	11
6.4	Batched feedback: the core reformulation	12
6.5	Memory as an explicit budget	14
6.6	Convolutional machines	14
6.7	Booleanization	15
6.8	Interpretation	16
6.9	Correctness	16
7	Usage	17
7.1	Thirty seconds	17
7.2	Tabular data, rules in the original units	17
7.3	Convolutional MNIST on a GPU	18
7.4	Explaining a single prediction	18
7.5	Building a new variant	19
8	Benchmarks	19
8.1	Method	19
8.2	Throughput against batch size	20
8.3	Throughput against model size	21
8.4	The model zoo	21
8.5	Where an update actually goes	22
8.6	CPU threads	22
8.7	Where the data lives	23
8.8	End to end	23
8.9	What batching costs in accuracy	23
9	Applications	24
9.1	What ships with the library	24
9.2	Where a Tsetlin machine earns its place	24
10	Limitations and future development	26
10.1	Known limitations	26
10.2	Where it goes next	27
11	Availability	28
11.1	Getting started	28
11.2	Reproducing the measurements	28
11.3	Quality gates	28
11.4	Documentation	29
11.5	Citation	29
11.6	Contributing	29
	References	29

1 Executive summary

How to read this document

Sections 1–4 state the problem and what the library offers; a reader who only wants to know *whether* to use torchsetlin can stop after Section 4. Sections 5–7 are the engineering core: architecture, the batched learning algorithm, and the API as it is actually used. Section 8 reports measured performance, and Sections 9–11 cover applications, known limits and where to get the code.

torchsetlin is an open-source Python package that implements the Tsetlin machine family as first-class PyTorch modules. A Tsetlin machine learns *propositional rules* — clauses such as `IF x0 AND NOT x1 THEN class 1` — using teams of Tsetlin automata driven by reward and penalty signals rather than gradients [1]. The resulting model is not merely explainable after the fact; the model *is* the explanation, and evaluating it costs a conjunction over included literals plus a vote count.

The problem torchsetlin addresses is not the algorithm, which is well established, but its software. Tsetlin machine implementations have historically been standalone C or CUDA kernels wrapped in scikit-learn-shaped Python classes. That design serves a benchmark script well and serves a research workflow badly: models do not move between devices the way the rest of a practitioner’s stack does, they do not serialise like other models, they do not consume a `DataLoader`, and extending them means writing a kernel rather than a few lines of tensor code.

torchsetlin takes the opposite starting point. The learnable state — one integer counter per (clause, literal) pair — is a `torch` buffer inside an `nn.Module`. Everything else follows from that decision:

- **One line to learn.** `model.update(x, y)` is the analogue of `loss.backward(); optimizer.step()`. It returns the pre-update vote sums, so a training loop gets batch metrics without a second forward pass.
- **One line to move.** `model.to("cuda")` moves the automata; there is no separate device build, no kernel recompilation and no host round-trip through NumPy.
- **Everything else is PyTorch.** `state_dict()`, `load_state_dict()`, `DataLoader`, `Dataset`, hooks, `torch.compile` and profiling all work because there is nothing non-standard to work around.
- **Batched feedback.** The learning update is reformulated so that a mini-batch is consumed in a handful of tensor operations rather than one stochastic pass per example. This is what makes GPU execution worthwhile (Section 6.4).
- **A complete workflow.** Booleanization encoders, a Keras-style `Trainer` with callbacks, metrics, rule extraction, closed-form feature importance and plotting ship in the same package.

At a glance

Table 1 is the whole package in one view.

Table 1. torchsetlin at a glance.

Models	<code>TsetlinMachine</code> (multi-class, optional integer clause weights), <code>CoalescedTsetlinMachine</code> (shared clause pool; multi-class or multi-label), <code>RegressionTsetlinMachine</code> , and convolutional variants <code>ConvTsetlinMachine</code> , <code>ConvCoalescedTsetlinMachine</code> , <code>ConvRegressionTsetlinMachine</code> , <code>Conv1dTsetlinMachine</code>
Learning	vote margin T , specificity s , memory depth N , boosted true-positive feedback, clause-size constraint, drop-clause/drop-literal, focused negative sampling, batched or exact sequential feedback
Data	thermometer, one-hot, bit-plane, adaptive-threshold, colour-thermometer and hypervector encoders; synthetic datasets; torchvision helpers
Training	<code>Trainer</code> over tensors/ <code>Dataset</code> / <code>DataLoader</code> ; early stopping, checkpoints, CSV logging, hyper-parameter schedules, <code>History</code>
Interpretation	rule extraction, clause activity and precision, closed-form global and local feature importance, per-example explanations
Peak measured	2.20×10^5 training examples/s and 1.05×10^6 inference examples/s (RTX 3090, 784 features, 500 clauses/class); up to 273× a 16-core CPU on the same code
Footprint	pure Python + PyTorch ≥ 2.0 and NumPy; no compiler, no CUDA toolkit, no C extension; Python 3.9–3.12; MIT licence

In short. A Tsetlin machine is a matrix of small integers updated by counting events. Once that is said in tensors instead of in C, the algorithm inherits the entire PyTorch ecosystem for free — and a GPU turns a 114-second MNIST epoch into a 1.5-second one.


2 Motivation

2.1 Why a logic learner at all

Deep networks answer; they rarely account for their answers. In clinical decision support, regulated industry and on-device inference, the account matters as much as the answer: a cardiologist wants to know which measurement moved the prediction, an auditor wants a decision rule they can read, and a microcontroller wants an inference path that is a handful of bit operations rather than a matrix cascade.

The Tsetlin machine ^[1] occupies exactly that niche. It learns conjunctive clauses over Boolean literals; clauses vote for or against each class; the class with the largest vote sum wins (Figure 1). Three properties follow:

- **Legibility by construction.** A clause is a readable conjunction. There is no surrogate model, no saliency heuristic, and no gap between what is explained and what is executed. Global and local importance have closed-form expressions ^[4].
- **Bit-level economy.** Inputs are Booleans, the learned state is small integers, and inference is an AND over included literals plus a vote count — which is why a substantial hardware literature has grown around it, from FPGA and ASIC inference engines to in-memory computing ^[15,16,17].



Every clause is a conjunction the reader can check by hand; the prediction is a signed vote count over the clauses that matched. The model is its own explanation, and inference is an AND plus an addition.

Figure 1. A Tsetlin machine classifier at inference time. Features are expanded into literals $[x, \neg x]$; each clause is a conjunction of the literals its automata have decided to *include*; matching clauses cast signed votes; the argmax over clipped vote sums is the prediction.

- **Competitive accuracy** on the tasks where those properties are worth paying for: tabular and clinical data, keyword spotting, and small-image classification.

2.2 Three obstacles in practice

The motivation for torchsetlin came from trying to use Tsetlin machines in a medical-imaging research group, where the surrounding stack is PyTorch and the data are large. Three obstacles showed up immediately, and none of them is about the algorithm.

1. The learning rule looks hostile to a GPU. The classical algorithm is defined per example: evaluate the clauses, draw feedback, *update the automata*, then move to the next example. The update is stochastic and touches every (clause, literal) counter, so a naive implementation performs one small, serial, random-number-heavy pass per training example. On a device with thousands of cores and a kernel-launch cost, that is the worst possible shape of work. A GPU implementation therefore requires reformulating the update, not merely porting it.

2. The tooling lives outside the ecosystem. Existing implementations are compiled kernels behind thin wrappers (Section 3). They are fast and faithful, but they sit beside a PyTorch pipeline rather than inside it. Moving a model to a GPU, checkpointing it mid-training, feeding it from a `DataLoader` that already augments on-device, or profiling it next to the CNN it is being compared against — each requires bespoke glue, and the data usually makes a round trip through host memory and NumPy on the way.

3. Everything around the model is re-invented per project. A Tsetlin machine consumes Booleans, so every project starts by writing a Booleanizer: thermometer thresholds for tabular columns, adaptive thresholding for images, bit planes for integers. Then a training loop, then early stopping, then a rule printer, then a clause visualiser. This is the unglamorous 80% of a Tsetlin machine project, it is rewritten constantly, and small differences in it (which thresholds, which threshold strategy) move accuracy more than the model hyper-parameters do.

The design question

Can the Tsetlin machine be expressed so that it is *ordinary* inside PyTorch — same device semantics, same checkpoints, same data plumbing — without giving up fidelity to the published algorithm? torchsetlin is the answer to that question: yes, if the learning update is reformulated from per-example state mutation into per-batch *event counting* (Section 6.4), and if the exact per-example path is kept available for when fidelity matters more than speed.

3 Existing solutions and what they leave open

3.1 The landscape

Tsetlin machine software has been developed almost entirely by the research groups that developed the algorithms, and it reflects that origin: each implementation is organised around a kernel that is fast and faithful to one paper, with a thin Python surface on top.

Reference CPU implementations. The canonical Python packages (`pyTsetlinMachine` and successors) implement the machine in C and expose it through a wrapper offering `fit/predict/transform`, one class per published variant (multi-class, convolutional, regression, weighted). They are the semantic ground truth for the algorithm and remain the reference any other implementation should be checked against.

CUDA implementations. `PyTsetlinMachineCUDA` and its relatives move the same algorithm into hand-written CUDA kernels, typically compiled at import time through `PyCUDA`. They demonstrated that the Tsetlin machine *can* be run at scale on a GPU — but as a separate package with a separate class hierarchy, so a project that needs both CPU and GPU maintains two code paths.

Unified architectures. More recent work (the Tsetlin Machine Unified project, `tmu`) factors the machine into reusable components — clause banks, weight banks — so that new variants can be assembled rather than rewritten, with CPU and optional CUDA backends behind one interface. This is the closest in spirit to `torchsetlin`: a shared substrate instead of a family of one-off programs. The substrate, however, is a custom C/CUDA layer rather than a tensor library, so the interoperability question below is unchanged.

Specialised and third-party work. Beyond these sit dedicated packages for particular variants (graph Tsetlin machines ^[14], sparse and coalesced CUDA machines), community ports in Rust, Julia, Go and JavaScript, and a growing hardware literature — FPGA and ASIC inference engines, asynchronous logic designs and in-memory computing architectures ^[15,16,17] — which target deployment rather than research iteration.

3.2 The gap

None of this is a criticism of the implementations on their own terms; they are fast, they are correct, and several of them predate the workflows described here. The gap is structural, and it has four parts.

1. **No shared tensor abstraction.** The learned state lives in a private C or CUDA buffer. It cannot be inspected, sliced, masked, moved or saved with the tools a PyTorch user already has; every such operation needs an accessor written for it.
2. **Device transfer is a packaging decision, not a runtime one.** Which device a model runs on is chosen by which package was imported, so CPU debugging and GPU training are different code paths rather than `.to("cpu")` and `.to("cuda")`.
3. **The data path leaves the framework.** Booleanized batches are NumPy arrays on the host. A pipeline that already has its data on the GPU — the normal case when the Tsetlin machine is being compared against, or combined with, a neural network — pays a round trip in both directions.

Table 2. Structural comparison of implementation approaches. ● available, ● partial or variant-specific, ○ not applicable to the design. The table compares *integration properties*, not accuracy or peak kernel speed; no head-to-head benchmark against other packages was run for this paper (Section 10).

Property	Reference C packages	CUDA packages	Unified C/CUDA	torchsetlin
Learned state is a framework tensor	○	○	○	●
<code>.to(device)</code> moves the model	○	○	●	●
Same code path on CPU and GPU	○	○	●	●
<code>state_dict()</code> checkpointing	○	○	○	●
Consumes <code>Dataset/DataLoader</code>	○	○	○	●
Accepts device-resident input	○	●	●	●
New variant without compiled code	○	○	●	●
Booleanization encoders included	●	●	●	●
Training loop with callbacks	○	○	○	●
Interpretability utilities	●	●	●	●
Plotting utilities	○	○	●	●
Requires a C/CUDA toolchain	●	●	●	○
Hand-tuned bit-packed kernels	●	●	●	○

4. **Extension means writing a kernel.** Trying a modified feedback rule, an alternative negative-sampling scheme or a new output layer requires C or CUDA, which puts the experiment out of reach of most of the people who want to run it.

Table 2 summarises this. The last row is deliberate and cuts the other way: a hand-written kernel can pack literals into machine words and exploit clause sparsity in ways a generic tensor library cannot, so for a small machine on a CPU the reference implementations remain the faster choice. `torchsetlin` trades that per-operation efficiency for parallel width and for integration — a trade that pays exactly when the machine is large and a GPU is present, and that is quantified in Section 8.

In short. The algorithms are well served; the *workflow* is not. What has been missing is a Tsetlin machine that behaves like every other model in a modern Python stack — which means building it on a tensor library rather than beside one.

4 The proposed library

4.1 Design principles

`torchsetlin` is built on five commitments, in priority order. Where they conflict, the earlier one wins.

1. **Be a `torch.nn.Module`, not a wrapper around one.** The automata state is a registered buffer. Every consequence of that — device moves, checkpoints, `repr`, hooks — is inherited rather than implemented.
2. **Be faithful to the published algorithm, and be explicit where you are not.** The exact

per-example algorithm is always available (`feedback_mode="sequential"`). The default batched path is an approximation in the same sense as mini-batch SGD, its fidelity is characterised empirically (Section 8.9), and the trade is documented rather than hidden.

3. **Keep the learning mathematics stateless and separate.** All of it lives in one module, `functional.py`, as pure tensor functions. Models are thin, stateful compositions of those functions, which is what makes a new variant a subclass rather than a fork.
4. **Ship the whole workflow.** Booleanization, training, evaluation, interpretation and visualisation are part of the package, because in practice they are part of the job.
5. **Depend on as little as possible.** PyTorch and NumPy are the only hard requirements. Matplotlib, torchvision and scikit-learn are optional and import-guarded. There is no build step.

4.2 What the library provides

Models. Seven concrete machines, all sharing one base class and one learning loop (Table 3). The flat and convolutional families are not parallel hierarchies: convolution is a mix-in that overrides input encoding, clause evaluation and feedback counting, and is combined with each flat model to produce its convolutional counterpart (Section 6.6).

Table 3. The model zoo. F is the number of Boolean features, C the total clause count, K the number of outputs. Every model stores its automata in one $(C, 2F)$ integer buffer.

Class	Task	Output layer and distinguishing feature
<code>TsetlinMachine</code>	multi-class	$C = K \cdot n$ clauses, half positive and half negative polarity per class; optional integer clause weights (weighted TM ^[3])
<code>CoalescedTsetlinMachine</code>	multi-class, multi-label	one shared clause pool with a signed integer weight per (clause, output); clauses may change sides during learning ^[5]
<code>RegressionTsetlinMachine</code>	regression	all clauses positive; vote sum in $[0, T]$ mapped linearly onto the target range ^[7]
<code>ConvTsetlinMachine</code>	image classification	clauses slide over $k_h \times k_w$ patches and fire if <i>any</i> patch matches; thermometer-encoded patch coordinates ^[2]
<code>ConvCoalescedTsetlinMachine</code>	image, multi-label	convolutional clauses over a shared pool
<code>ConvRegressionTsetlinMachine</code>	image regression	reads a continuous quantity off an image ^[7]
<code>Conv1dTsetlinMachine</code>	sequences, signals	$1 \times k$ windows over (B, Z, L) Boolean channels

Learning options. The switches that the literature has shown to matter are constructor arguments rather than forks: vote margin T and specificity s ; memory depth N ; boosted true-positive feedback; integer clause weights; the clause-size constraint of ^[6]; drop-clause and drop-literal regularisation ^[8]; focused negative sampling; and the choice between batched and exact sequential feedback.

Booleanization. Seven encoders — a plain binarizer, thermometer (quantile, uniform or unique-value thresholds), one-hot, bit plane, adaptive Gaussian/mean threshold, colour thermometer and sparse hypervector ^[13] — all implemented as `nn.Modules` with the scikit-learn `fit/transform` protocol, composable with `Compose`, movable to a device and saved in a `state_dict`. Encoders also carry their inverse knowledge: `ThermometerEncoder.decode_range` turns included bits back into an

interval in the original units, so a rule reads `IF age>=52.4 AND NOT age>=67.1` rather than `IF x37 AND NOT x41`.

Training. A `Trainer` with `fit/evaluate/predict` over tensors, `Datasets` or `DataLoaders`; a `History` object; and callbacks for early stopping, checkpointing, CSV logging, hyper-parameter scheduling and automata-state logging. Nothing in the models requires the trainer — a hand-written loop over `model.update` is a first-class way to use the library.

Evaluation and interpretation. Accuracy, confusion matrix, per-class precision/recall/F1, regression and multi-label metrics, expected calibration error and trustworthiness curves; rule extraction, clause activity and precision, closed-form global and local feature importance [4], and per-example explanations listing the clauses that fired and how they voted.

Visualisation. Training curves, confusion matrices, automata-state heat maps, clause memory plots, convolutional clause patches rendered as pixel stencils, vote distributions and trustworthiness plots.

What torchsetlin is not

It is not a reimplement of every Tsetlin machine variant in the literature — graph machines, autoencoders, absorbing automata and TM composites are not implemented (Section 10). It is not a bit-packed kernel, and on a small machine on a CPU it will not beat one. It is the substrate on which those things can be built in tensor code instead of C.

5 System architecture

5.1 Four layers

`torchsetlin` is organised as four layers with a strict one-way dependency (Figure 2). Each layer is useful on its own: a researcher testing a new feedback rule works in layer 1, someone adding a model variant works in layer 3, and an applied user only ever sees layer 4.

Layer 1 — `functional.py`. Stateless tensor functions: literal expansion, clause evaluation, vote sums, feedback probabilities, negative-class sampling, Type I and Type II event counting, and the automaton transition itself. No state, no autograd, no device assumptions. All learning mathematics lives here, which means it can be unit-tested against a brute-force reference implementation and reused by code that does not use the model classes at all.

Layer 2 — `TsetlinMachineBase`. An `nn.Module` that owns the automata and runs the generic learning loop: input coercion, lazy shape initialisation, chunking, dropout masks, feedback accumulation and the commit. It knows nothing about classes, labels or images.

Layer 3 — concrete models. Each variant supplies only an *output layer* (how clause outputs become votes) and a *feedback policy* (which clauses receive which feedback for a labelled example). Convolution is orthogonal to both and enters as a mix-in, so `class ConvTsetlinMachine(_ConvMixin, TsetlinMachine)` is the entire definition of the convolutional classifier’s learning behaviour.




Figure 2. The four layers. Dependencies point downwards only: the learning mathematics does not know what a model is, and a model does not know what a trainer is.

Layer 4 — workflow. Data preparation, training, metrics, interpretation and plotting. These depend on the models through their public API only, so a user’s own model subclass gets the whole workflow layer for free.

5.2 The learning state

The learnable state is a **buffer**, not a **parameter**, because a Tsetlin machine has no gradients:

Buffer	Shape / dtype	Role
<code>ta_state</code>	$(C, 2F)$ <code>int32</code>	the automata; literal included iff state $\geq N$
<code>include</code>	$(C, 2F)$ <code>float</code>	cached 0/1 mask, derived, non-persistent
<code>include_count</code>	$(C,)$ <code>long</code>	literals per clause, derived, non-persistent
<code>weights</code>	$(C,)$ or (C, K) <code>int32</code>	clause weights (weighted/coalesced models)

Only `ta_state` and `weights` are persistent; the two caches are recomputed whenever the state changes, including after `load_state_dict`. Storing the derived mask costs $2CF$ floats and removes a comparison from the inner loop of every forward pass, which is the single hottest path in the library.

Two consequences are worth calling out. First, `model.parameters()` is deliberately empty and warns when it returns nothing, because silently returning an empty iterator to someone who wrote

`optim.SGD(model.parameters())` would be worse than an error. Second, `n_features=None` enables lazy initialisation in the style of `torch.nn.LazyLinear`: the state is allocated on the first batch, and `_load_from_state_dict` reconstructs the shape from the checkpoint before loading so that a lazily built model can be restored without ever having seen data.

5.3 How an update flows

`update(x, y)` is the analogue of `loss.backward()`; `optimizer.step()`, and it returns the *pre-update* vote sums so that a training loop can log batch accuracy without a second forward pass. Figure 3 shows the flow.




Figure 3. The `update()` pipeline. The batch is split into chunks that respect a memory budget; each chunk contributes *event counts* to a shared accumulator; the automata are written once, at the commit. Amber marks the hook a new model must implement.

The separation between *accumulate* and *commit* is the structural core of the library. Accumulation is pure counting and can be done in any number of chunks, in any order, without changing the result. The commit is the only place the automata are written, the only place randomness enters the state, and the only place the derived caches are refreshed. That split is what makes the batched algorithm expressible, what bounds memory usage, and what makes the convolutional mix-in possible — convolution changes what is counted and how, but not when or how it is committed.

5.4 Extension points

A new model is a subclass, not a fork. The required surface is three methods:

Method	Contract
<code>_votes(clause_out)</code>	(B, C) Boolean clause outputs $\rightarrow (B, K)$ vote sums
<code>_select_feedback(votes, y, clause_out)</code>	$\rightarrow$ two (B, C) Boolean masks: who gets Type I, who gets Type II
<code>_coerce_targets(y, n, device)</code>	validate and convert the targets
<i>optional</i>	
<code>_accumulate_weights / _commit_weights</code>	clause weights
<code>_encode / _evaluate / _feedback_counts</code>	a different input structure (this is the convolutional mix-in)
<code>_chunk_elements_per_example</code>	the memory model for that structure
<code>rules()</code>	how the model prints itself

In short. Learning mathematics is stateless and lives in one file; state and the loop live in one base class; variants supply an output layer and a feedback policy. The entire convolutional family is one mix-in over that structure — roughly 360 lines, none of which re-implement learning.

6 Technical design

This section states the algorithm as torchsetlin implements it. The notation throughout: B examples per batch, F Boolean features, $2F$ literals, C clauses, K outputs, N automaton states per action, s specificity, T vote margin.

6.1 Representation

An example is a Boolean vector $x \in \{0, 1\}^F$. Its *literal vector* is $L(x) = [x, \neg x] \in \{0, 1\}^{2F}$. Every (clause, literal) pair owns a Tsetlin automaton whose state $\theta_{jk} \in \{0, \dots, 2N - 1\}$ decides membership:

$$\text{include}_{jk} = \mathbb{1}[\theta_{jk} \geq N], \quad c_j(x) = \bigwedge_{k: \text{include}_{jk}} L_k(x).$$

The whole learned state is therefore one integer matrix $\theta \in \mathbb{Z}^{C \times 2F}$ — in torchsetlin, the `int32` buffer `ta_state`. Two derived, non-persistent buffers cache the include mask as floats and the per-clause literal count; anything that mutates θ refreshes them, including `load_state_dict`.

6.2 Inference is two matrix products

A conjunction is awkward to vectorise directly, but its negation is a sum. A clause fails on an example exactly when it includes at least one literal that is False, so

$$\text{violations} = (1 - L) \text{include}^\top \in \mathbb{R}^{B \times C}, \quad c_j(x_b) = \mathbb{1}[\text{violations}_{bj} = 0].$$




Figure 4. One Tsetlin automaton. The state line is split at N : below it the literal is forgotten, at or above it the literal is included in the clause. Deep states are hard to dislodge, states near the boundary flip easily — this is the memory that makes the machine robust to noisy examples.

One dense matrix product evaluates every clause on every example. The vote sums are a second product, $v = cW$, clipped to $[-T, T]$, where $W \in \mathbb{Z}^{C \times K}$ holds the signed clause weights. Prediction is $\hat{y} = \arg \max_k v_k$.

Empty clauses are mode-dependent — and this is not a detail

A clause with no included literals has zero violations, so the formula above makes it match *everything*. That is correct while learning: a fresh clause must fire in order to receive the Type Ia feedback that gives it its first literal. It is wrong at prediction time, where an empty clause should vote for nothing. torchsetlin keys this off `self.training`, exactly like `nn.Dropout`, so `model.eval()` is required before measuring accuracy — forgetting it silently changes predictions rather than raising. The behaviour is pinned by tests.

6.3 The feedback policy

For a labelled example the machine decides, per clause, whether to give Type I feedback (make the clause recognise this example), Type II feedback (make it stop matching), or nothing. The decision is probabilistic in the vote margin: the target class receives feedback with probability $(T - v_y)/2T$ and a sampled negative class with probability $(T + v_q)/2T$. Once a class already wins by T votes its clauses stop being updated, which frees them to specialise elsewhere — the resource-allocation mechanism at the heart of the algorithm [1].

Each model differs only in how votes are formed and who receives which feedback (Table 4). Everything below that line is shared.

Two refinements are available on the classifiers. *Integer clause weights* (the weighted TM [3]) let one clause count for many, which is what makes large convolutional machines tractable: Type Ia

Table 4. Output layer and feedback policy per model. This table is the difference between the model classes; the learning machinery underneath is identical.

Model	Vote sums	Who receives what
TsetlinMachine	$v_k = \sum_{j \in \text{pos}(k)} w_j c_j - \sum_{j \in \text{neg}(k)} w_j c_j$	target class: Type I to positive, Type II to negative clauses; one random negative class: the reverse
CoalescedTsetlinMachine	$v = c W, W \in \mathbb{Z}^{C \times K}$	per output, Type I when the weight sign agrees with the target and Type II when it does not; weights move by ± 1 , so a clause can change sides
RegressionTsetlinMachine	$v = \sum_j w_j c_j$ clipped to $[0, T]$, mapped to $[y_{\min}, y_{\max}]$	Type I with probability $(v^* - v)/T$ when the prediction is too low, Type II with $(v - v^*)/T$ when too high

increments a clause’s weight, Type II decrements it. *Focused negative sampling* draws the negative class in proportion to how strongly it currently votes rather than uniformly, so the machine corrects the classes it actually confuses first — worth having when K is large.

6.4 Batched feedback: the core reformulation

The classical algorithm applies feedback one example at a time, re-evaluating the clauses in between. Each step draws random numbers for the whole $(C, 2F)$ state and writes it back. For a machine with 5 000 clauses over 784 features that is a 7.8-million-element stochastic pass *per training example* — on a GPU, a sequence of small kernels with nothing to parallelise across.

torchsetlin reformulates the update around a simple observation: *every stochastic decision in Type I feedback is an independent Bernoulli trial with a state-independent probability*. Memorisation happens with probability $p_{\text{mem}} = (s - 1)/s$ (or 1 with boosted true positives), forgetting with probability $1/s$. Whether a particular (clause, literal) automaton is offered such a trial depends only on which feedback events the batch produced — not on the automaton’s own state. So the trials can be *counted* first and *sampled* once.

Let $S_I, S_{II} \in \{0, 1\}^{B \times C}$ mark which (example, clause) pairs receive Type I and Type II feedback among the clauses that *matched*, and S_I^{nofire} those that received Type I without matching. Then, with $L \in \{0, 1\}^{B \times 2F}$ the batch literal matrix:

$$n^{\text{true}} = S_I^T L, \quad n^{\text{false}} = S_I^T (1 - L), \quad (1)$$


$$n_j^{\text{lb}} = \sum_b S_I^{\text{nofire}} [b, j], \quad n^{\text{II}} = S_{II}^T (1 - L). \quad (2)$$

Four count tensors, obtained from two matrix products and a column sum. The transition then follows in closed form:

$$\Delta = \underbrace{\text{Bin}(n^{\text{true}}, p_{\text{mem}})}_{\text{Type Ia memorise}} - \underbrace{\text{Bin}(n^{\text{false}} + n^{\text{lb}}, \frac{1}{s})}_{\text{Type Ia forget + Type Ib}} + \underbrace{\min(n^{\text{II}}, N - \theta)}_{\text{Type II, bounded at the boundary}}, \quad (3)$$


and $\theta \leftarrow \text{clip}(\theta + \Delta, 0, 2N - 1)$. The two forgetting sources merge because both are Bernoulli($1/s$) trials and a sum of independent Bernoulli trials with equal probability is binomial. Type II is deterministic, and the min reproduces a property of the sequential algorithm exactly: an excluded False literal is pushed towards inclusion one step at a time, but once it crosses the boundary the clause no longer matches and receives no further Type II events in that update.

A Classical sequential algorithm — B stochastic commits per mini-batch



Each example sees the automata the previous example left behind. That serial dependency is the *point* of the algorithm — and the reason a GPU sits idle: the work between two state writes is far too small to fill it.

B Batched formulation — count the events, then commit *once*



Every stochastic decision of the sequential algorithm is an independent Bernoulli trial, so the trials one automaton receives in a batch sum to a *binomial* draw — one $(C, 2F)$ random tensor per batch instead of one per example. At $B = 1$ this reduces exactly to the classical algorithm.

Figure 5. The reformulation. The sequential algorithm (A) performs one stochastic $(C, 2F)$ pass per example. The batched algorithm (B) accumulates the same events as counts and realises them in a single binomial draw per batch. Cost per update becomes almost independent of B , which is exactly the property a GPU rewards.

What is exact and what is not. At $B = 1$ Equation (3) reduces to the classical algorithm term by term. For $B > 1$ it is an approximation in precisely the sense mini-batch SGD approximates SGD: the examples within a batch do not see each other's updates. Two secondary effects follow — an automaton can move by more than one step in a commit, and the Type II bound is evaluated against the state at commit time rather than after each event. Section 8.9 measures what this costs. For work where fidelity matters more than throughput, `feedback_mode="sequential"` (or `update(..., sequential=True)`) runs the exact per-example algorithm, still vectorised across clauses and literals.

A design decision that measurement reversed

An obvious refinement is to *thin* the accumulated events — to cap how many Type Ia or Type II events one clause may contribute per batch, on the grounds that the sequential algorithm would have re-evaluated in between and delivered fewer. This was implemented and measured; it made learning distinctly worse. Aggregating all events is therefore deliberate, not an oversight, and is documented as such in the repository so that the next person to have the idea finds the answer before spending a week on it.

6.5 Memory as an explicit budget


Intermediates, not the model, dominate memory. `torchsetlin` never silently allocates a tensor proportional to B without a bound: each model declares how many elements one example costs,

$$\text{flat: } 4(C + 2F), \quad \text{convolutional: } 3PC + 2C \cdot 2F + P \cdot 2F,$$

where P is the number of patches per image, and the base class splits the batch into chunks that keep the largest intermediate under `max_chunk_elements` (default 2^{27} elements). Chunking never changes the result of a batched update — accumulation is pure addition — so it is purely a memory-for-nothing trade. Getting the per-example estimate wrong is not benign, however: an under-estimate produces many tiny chunks per batch and a 4–20× slowdown, which is why the estimate is part of the subclass contract rather than a heuristic.

6.6 Convolutional machines

A convolutional Tsetlin machine [2] lets a clause look for a pattern *anywhere* in an image instead of memorising absolute positions. Each $k_h \times k_w$ window becomes a patch whose Boolean features are the $Zk_h k_w$ pixels followed — when `position_encoding=True` — by thermometer-encoded window coordinates, $P_y - 1$ bits [$y > k$] and $P_x - 1$ bits [$x > k$]. A clause is True for the image if it matches at least one patch, and a feedback event updates the clause from *one uniformly random matching patch*.



Drawing the patch here rather than during evaluation keeps *prediction* free of randomness — a (B, P, C) random tensor on every forward pass cost roughly 5× the runtime — and lets Type Ia and Type II draw independently, which matters for coalesced and multi-label models where one clause can receive both in the same update.

Figure 6. Convolutional evaluation. `unfold` produces the patch literals, one matmul produces all (patch, clause) matches, and an OR over patches gives the clause output. The random matching patch is drawn during *feedback counting*, over selected events only.

Implementation-wise this is `torch.nn.functional.unfold`, one matmul, and an `any` reduction. The subtlety is *where* the random patch is drawn. Drawing it during evaluation is the obvious choice and it is wrong twice over: it costs a (B, P, C) random tensor and an `argmax` on *every* forward pass, including pure prediction where the draw is discarded — roughly 5× the runtime on MNIST-sized

inputs — and it forces Type Ia and Type II to share one patch, which is incorrect for coalesced and multi-label models where a clause can receive both in the same update. `torchsetlin` therefore returns the per-patch match tensor as an opaque *feedback context* from evaluation and draws inside the feedback counter, over the selected events only. Prediction is RNG-free, and a test pins that it stays so.

The random draw itself is a small trick worth noting: for candidate mask $\text{cand} \in \{0,1\}^{n \times P}$, `argmax(rand(n,P) * cand)` selects a uniformly random index among the ones, with no rejection loop and no gather over a variable-length list.

Position encoding is a modelling choice, not a default to leave alone

Position literals let a clause demand *where* it matches. On genuinely translation-invariant tasks they hurt — on the generated `make_shapes` benchmark, accuracy oscillates between 0.5 and 0.9 with them enabled and reaches 1.000 on every seed with them disabled. On MNIST, where digits are centred and position is informative, they are worth about a point of accuracy. Check this flag before suspecting the feedback code.

6.7 Booleanization

The literals are the vocabulary of every rule the machine can learn, so Booleanization is the most consequential modelling decision in a Tsetlin machine project — and in practice it moves accuracy more than the model hyper-parameters do. `torchsetlin` treats encoders as first-class modules (Table 5): they are `nn.Modules` with a `fit/transform` protocol, they move to a device, they are saved in a `state_dict`, and they compose.

Table 5. Booleanization encoders. F input columns, b bits per column.

Encoder	Output	Use and why
<code>Binarizer</code>	F	a fixed threshold per element, the original MNIST recipe (<code>pixel > 0.3</code>)
<code>ThermometerEncoder</code>	$F \cdot b$	ordinal bits $[v \geq t_1, \dots, v \geq t_b]$; two literals then express an interval (<code>age >= 30 AND NOT age >= 50</code>). Thresholds by quantile, uniform width or unique value
<code>OneHotEncoder</code>	$\sum_i n_i$	one bit per category; rules read as <code>colour=red</code>
<code>BitPlaneEncoder</code>	$F \cdot b$	place-value bits of integers; compact codes
<code>AdaptiveThresholdEncoder</code>	F	local Gaussian/mean thresholding, the standard recipe for MNIST-like images in the convolutional TM literature
<code>ColorThermometerEncoder</code>	$C \cdot b$ planes	per-channel thermometer for RGB
<code>HypervectorEncoder</code>	d	sparse random hypervectors bundled with OR, optionally binding position by cyclic shift — a compact encoding for large vocabularies ^[13]

Encoders also carry the inverse mapping, which is what keeps rules readable. A thermometer encoder’s `decode_range` turns a clause’s included bits back into an interval in the original units, and `aggregate_literal_importance` sums literal importances back over the bits of one original column.

A GPU bug worth documenting

Encoders keep their thresholds on the CPU, so handing one a CUDA tensor triggers a device-to-host copy. Issued with `non_blocking=True` into pageable memory, such a copy returns *before* the data lands, and the encoder thresholds stale memory — about 40% wrong bits, no error, just a quietly worse model. All coercion in `torchsetlin` routes through a helper that sets `non_blocking` only for CUDA destinations. The regression test has to allocate on the device (`torch.rand(..., device=dev)`) rather than moving a host tensor across, or the race never occurs and the test passes vacuously.

6.8 Interpretation

Three levels of explanation are available, all closed-form and all computed from the state rather than by probing the model.

Rules. `model.rules()` prints each non-empty clause as `IF <conjunction> THEN <class>`, with its weight where weights exist. Feature names flow through from the encoder, so the conjunction is in the user’s vocabulary.

Global importance. Following ^[4], the strength of literal l for output k is the weighted fraction of positive clauses of k that include it,

$$g_k[l] = \frac{1}{m_k} \sum_{j: W_{jk} > 0} W_{jk} \text{include}_{jl}, \quad m_k = |\{j : W_{jk} > 0\}|.$$

Local importance and explanations. For one example, the same quantity restricted to the clauses that actually matched and to the literals that are True gives a per-example attribution; `interpret.explain` returns the matching clauses ranked by weight, which is an account of the prediction in the model’s own terms rather than a surrogate.

6.9 Correctness

The library is tested with 44 tests, parametrised over a `device` fixture so that every device-marked test runs twice on a GPU machine. The tests that matter most are not the smoke tests but the semantic pins: brute-force reference implementations of clause evaluation and of convolutional patch semantics; the empty-clause train/eval asymmetry; the RNG-free prediction path for convolutional models; the encoder device-transfer race described above; and a packaging test asserting that the version agrees between `pyproject.toml` and `__init__.py` and that every subpackage ships in the wheel — which exists because a `.gitignore` rule once silently excluded `torchsetlin/data` from a release.

In short. Clause evaluation is a matmul against the include mask; learning is event counting plus one binomial draw per commit. Those two identities are the whole library: everything else — models, convolution, chunking, the trainer — is organisation around them.

7 Usage

The API has two levels: a bare loop over `model.update` for people who want control, and a `Trainer` for people who want Keras. Both are shown below; every listing is runnable as written, and all of them have full counterparts in the repository's `examples/` directory.

7.1 Thirty seconds

```
import torchsetlin as tt

x_train, y_train = tt.data.make_noisy_xor(5000, noise=0.4) # Boolean features, int labels
x_test, y_test = tt.data.make_noisy_xor(2000, noise=0.0, seed=1)

model = tt.TsetlinMachine(n_features=12, n_classes=2, n_clauses=20, T=15, s=3.9).to("cuda")

# update() is the analogue of loss.backward(); optimizer.step()
for epoch in range(50):
    model.train()
    for i in range(0, 5000, 10):
        model.update(x_train[i:i+10].cuda(), y_train[i:i+10].cuda())

model.eval() # empty clauses now vote for nothing
acc = (model(x_test.cuda()).argmax(1) == y_test.cuda()).float().mean()
print(model.rules()[:2])
# ['IF x0 AND NOT x1 THEN 1', 'IF x1 AND NOT x0 THEN 1']
```

The last line is the point of the exercise: the model has recovered the generating rule of the Noisy XOR benchmark from data with 40% label noise, and printed it.

7.2 Tabular data, rules in the original units

Continuous columns become Boolean through a thermometer encoder. Passing the encoder's feature names to the model is what turns `x37` back into `worst_perimeter>=105.9`, and `max_included_literals` keeps the printed rules short enough to read.

```
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
import torch, torchsetlin as tt

tt.seed_everything(0)
data = load_breast_cancer()
x_tr, x_te, y_tr, y_te = train_test_split(data.data, data.target, test_size=0.3, random_state=0)
x_tr, x_te = torch.tensor(x_tr).float(), torch.tensor(x_te).float()

enc = tt.data.ThermometerEncoder(n_bits=8, strategy="quantile").fit(x_tr)
xb_tr, xb_te = enc(x_tr), enc(x_te)

model = tt.TsetlinMachine(xb_tr.shape[1], 2, n_clauses=100, T=20, s=4.0,
                          max_included_literals=6)
model.feature_names = enc.feature_names(list(data.feature_names)) # readable literals
model.class_names = list(data.target_names)

trainer = tt.Trainer(model, batch_size=8, callbacks=[tt.EarlyStopping(patience=30)])
trainer.fit((xb_tr, y_tr), epochs=200, val_data=(xb_te, y_te))
print(trainer.test((xb_te, y_te))) # {'test_accuracy': 0.96...}

for rule in model.rules(polarity=1)[:5]: # clauses voting *for* a class
    print(rule)
```

```
# Which original columns carry the model, with the thermometer bits summed back together
imp = tt.interpret.global_feature_importance(model, class_index=0)
per_feature = tt.interpret.aggregate_literal_importance(
    imp, [i // 8 for i in range(xb_tr.shape[1])])
```

7.3 Convolutional MNIST on a GPU

The convolutional path differs in three places only: the encoder (adaptive thresholding rather than a fixed cut), the model class, and the patch geometry. Keeping the encoder on the compute device means the images never travel back to the host.

```
import torch, torchsetlin as tt
dev = "cuda"

enc = tt.data.AdaptiveThresholdEncoder(block_size=11, C=2, method="gaussian").to(dev)
booleanize = lambda x: enc(x * 255.0)

x_train, y_train = tt.data.load_mnist_boolean("./data", train=True,
                                              encoder=booleanize, device=dev)
x_test, y_test = tt.data.load_mnist_boolean("./data", train=False,
                                             encoder=booleanize, device=dev)

model = tt.ConvTsetlinMachine(
    n_classes=10, n_clauses=200, T=100, s=10.0,
    patch_size=10, stride=1, position_encoding=True, weighted=True,
    input_shape=tuple(x_train.shape[1:]),          # (1, 28, 28)
)

trainer = tt.Trainer(model, device=dev, batch_size=32, callbacks=[
    tt.EarlyStopping(monitor="val_accuracy", patience=5),
    tt.ModelCheckpoint("mnist_conv.pt", monitor="val_accuracy"),
    tt.CSVLogger("mnist_conv.csv"),
    tt.StateSummaryLogger(),                        # mean literals/clause, empty clauses
])
trainer.fit((x_train, y_train), epochs=10, val_data=(x_test, y_test))

votes = trainer.predict((x_test, y_test), return_votes=True)
print(tt.metrics.classification_report(votes, y_test))

tt.viz.plot_conv_clauses(model, range(16))          # what each clause looks for
```

Hyper-parameter scheduling

Tsetlin hyper-parameters are ordinary Python attributes, so annealing one is a callback and not a framework feature: `tt.HyperparameterSchedule("s", lambda epoch: 10.0 - 5.0 * epoch / 30)`, or the same for `drop_clause_p`. `model.state_summary()` reports mean included literals per clause and the number of empty clauses — the two numbers that tell you whether `s` is in the right range.

7.4 Explaining a single prediction

```
expl = tt.interpret.explain(model, x_test[0], feature_names=model.feature_names)
print(expl)                                # shape of the output; the numbers depend on the run
# prediction: 1 votes: [-14.0, 17.0]
# [+3 -> output 1] clause 42: IF worst perimeter>=105.9 AND NOT mean texture>=21.8
# [+2 -> output 1] clause 77: IF worst concave points>=0.15
# ...
```

```
torch.save(model.state_dict(), "model.pt")          # ordinary PyTorch checkpointing
reloaded = tt.TsetlinMachine(None, 2, n_clauses=100, T=20, s=4.0)  # lazy: shape from ckpt
reloaded.load_state_dict(torch.load("model.pt"))
```

The explanation is not a surrogate: those clauses are the model, and the listed votes add up to the reported vote sums.

7.5 Building a new variant

A new model supplies an output layer and a feedback policy; the base class provides everything else, including chunking, dropout, the batched commit and the whole workflow layer.

```
class MyTsetlinMachine(tt.TsetlinMachineBase):
    def __init__(self, n_features, n_clauses, T, **kw):
        self.T = T
        super().__init__(n_features, n_clauses, **kw)

    def _votes(self, clause_out, clamp=True):      # (B, C) bool -> (B, n_outputs)
        ...

    def _coerce_targets(self, y, n, device):      # validate / convert targets
        ...

    def _select_feedback(self, votes, y, clause_out): # -> (type_i, type_ii, aux) masks
        ...
```

Or skip the class hierarchy entirely and compose the functional layer directly — useful when prototyping a different feedback rule:

```
from torchsetlin import functional as F

lit      = F.to_literals(x)                      # (B, 2F)
include  = (state >= N).float()
c        = F.clause_outputs(lit, include, empty_value=True) # (B, C)
v        = F.vote_sums(c, weights, T)
n_true, n_false, n_ib = F.type_i_counts(lit, sel_fire_i, sel_nofire_i)
n2       = F.type_ii_counts(lit, sel_fire_ii)
F.apply_feedback(state, n_true, n_false, n_ib, n2, n_states=N, s=s)
```

It is plain tensor code, so `torch.compile`, a custom CUDA kernel or an alternative sampling scheme can be dropped in without touching anything else.

8 Benchmarks

8.1 Method

Every number in this section comes from one run of the benchmark driver shipped with the repository, `benchmarks/run_benchmarks.sh`, on the machine in Table 6. Each point is the median of three timed blocks of steady-state calls after a warm-up, with the device synchronised once per block. The micro-benchmarks use synthetic Boolean data at 20% density, resident on the device unless stated otherwise. The driver writes a JSON record per measurement and regenerates the figures and tables from it, so the whole section can be reproduced — or replaced with another machine’s numbers — with one command.

Table 6. Measurement environment.

CPU	AMD Ryzen 9 3950X, 16 cores / 32 threads (PyTorch using 16 threads)
GPU	NVIDIA GeForce RTX 3090 (24 GB)
Software	PyTorch 2.10.0+cu128 / CUDA 12.8, Python 3.10.11, torchsetlin 0.1.2
Reference config	784 Boolean features, 500 clauses per class, 10 classes, unless stated otherwise

What this section does and does not measure

It measures *this* library on two devices: the same Python code, the same algorithm, one run on a CPU and one on a GPU. It is not a comparison against other Tsetlin machine implementations — no such comparison was run — and it is not a state-of-the-art accuracy claim. Where accuracy appears it is there to show that the batched update learns what it should, not to compete with 400-epoch results from the literature.

8.2 Throughput against batch size

One `update()` call touches every automaton in the machine, so its cost depends far more on the *size of the machine* than on how many examples are in the batch. That is the whole point of the reformulation in Section 6.4, and Figure 7 is what it buys: on the GPU, training throughput rises from 435 examples/s at batch 1 to 219,850 at batch 4096 — a factor of 505 for feeding the same kernel more work.




Figure 7. Throughput against batch size for a 784-feature, 500-clause-per-class machine (log-log). Training on the GPU scales almost linearly with batch size until the device saturates. The dashed line is the counter-example: a 12-feature, 10-clause machine is less work than a kernel launch, and only overtakes the CPU past batch 1024.

At batch 256 the same model trains at 88,908 examples/s on the GPU against 408 on the CPU — 218×. Inference behaves differently: the CPU forward pass is already cheap, so the gap there tops out around 28×, reaching 1.05×10^6 examples/s on the GPU.

The dashed series is the honest counterpoint. A Noisy-XOR-scale machine (12 features, 10 clauses per class) is *slower* on the GPU than on the CPU for every batch size below 1024. Small machines belong on a CPU, and torchsetlin makes that a one-word change rather than a different package.

8.3 Throughput against model size

Both axes of model size cost the same thing — an update is $O(C \cdot 2F)$ element-wise work — so both throughput curves fall roughly as $1/x$. The interesting quantity is the *ratio*, in Figure 8: the CPU falls off immediately while the GPU stays flat until the machine is large enough to fill it, so the advantage *widens* with model size, from 51× at 50 clauses per class to 273× at 4000, and from 71× at 128 features to 286× at 8192.




Figure 8. The GPU advantage grows with the model. Big machines — the ones the Tsetlin machine literature actually uses for images — are exactly the ones worth moving to a GPU.

This matters for a practical reason. The published configurations that reach competitive accuracy on images use thousands of clauses per class; on a CPU those are overnight runs, and the resulting iteration speed is what keeps people away. At 2000 clauses per class the GPU trains at 25,934 examples/s against 96 on the CPU.

8.4 The model zoo

Weighted and coalesced variants cost about what the plain machine costs — they change the output layer and the feedback policy, not the per-clause work that dominates (Table 7). The convolutional models are the outlier at roughly 4× the cost per example, which is arithmetic rather than overhead: a 10×10 window over a 28×28 image is 361 patches, and every patch is evaluated against every clause.

Table 7. All models, 784 features, batch 128. Throughput in examples/s.

Model	Training			Inference		
	CPU	GPU	speedup	CPU	GPU	speedup
TsetlinMachine	268	49,638	185×	30,719	526,882	17×
TsetlinMachine (weighted)	268	49,475	185×	32,227	491,498	15×
CoalescedTsetlinMachine	178	43,579	244×	34,222	526,083	15×
RegressionTsetlinMachine	214	56,822	266×	44,567	577,994	13×
ConvTsetlinMachine	137	13,059	95×	140	17,942	128×
ConvCoalescedTsetlinMachine	132	13,137	100×	138	17,826	129×

GPU peak memory is 262 MB for the flat models and 298 MB for the convolutional ones at this configuration — small enough that clause budget, not memory, is the binding constraint on most cards.

8.5 Where an update actually goes

Figure 9 breaks one `update()` call into stages. The result is unambiguous: `apply_feedback` — the stochastic increment and decrement of the whole ($C, 2F$) state — is 92% of the 623 ms CPU call and 42% of the 2.8 ms GPU call. A standalone `torch.binomial` draw of the same shape costs 266.8 ms on the CPU against 0.22 ms on the GPU, making it the single most expensive operation in the commit.




Figure 9. Stage timings for one update (5000 clauses, 784 features, batch 256). Note the log scale: the CPU total is 623.4 ms and the GPU total 2.8 ms. Everything before the commit — encoding, clause evaluation, votes and the feedback-count matmuls — is comparatively free.

This is also the quantitative argument for batching. A batched update draws once per commit; a sequential one draws once per example. Whatever the batch size, the expensive part of the call happens exactly once.

8.6 CPU threads

More CPU threads buy surprisingly little in training — 1.4× from 1 to 32 threads (Figure 10). The batched update is a chain of large element-wise operations over the state, and the binomial draw that dominates it does not thread. Inference scales much better, 9.6× over the same range, because the forward pass is one large matmul-shaped reduction. If a CPU run is too slow, the lever is the clause budget or the batch size, not `torch.set_num_threads`.




Figure 10. CPU thread scaling (5000 clauses, 784 features, batch 256).

8.7 Where the data lives

Keeping the Boolean dataset in host memory and copying each mini-batch across costs at most 2% of throughput on this machine — the update is long enough to hide the transfer. Streaming a dataset that does not fit in VRAM is therefore free in practice. The copy that *does* bite is the one in the other direction, into a CPU-resident encoder, which is the hazard described in Section 6.

8.8 End to end

Micro-benchmarks flatter whichever device wins the micro-benchmark, so Table 8 reports whole training runs with `Trainer.fit` included — data loading, epoch loop, validation and all.

Table 8. End-to-end MNIST runs. Accuracy matches across devices to within run-to-run noise, as it should: only the wall clock moves.

Model	Device	train ex.	epochs	s / epoch	test acc.
TsetlinMachine 500/class	CPU	60,000	3	114.35	0.9542
	GPU	60,000	3	1.52	0.9548
ConvTsetlinMachine 200/class, 10×10	CPU	10,000	1	70.19	0.8270
	GPU	10,000	1	0.82	0.8240

A 114-second epoch becomes a 1.5-second one (75×); the convolutional run goes from 70 seconds to 0.8 (86×). That is the difference between a hyper-parameter sweep you plan and one you simply run.

On accuracy: the flat 500-clause machine reaches 95.8% after five epochs, and the convolutional model 97.2% after three epochs at batch 32 (200 clauses per class, $T = 100$, $s = 10$, weighted). The published figures for much larger budgets — around 98% for 2000 clauses per class after 400 epochs, and around 99% for convolutional machines with 8000 clauses per class — are what the same code should reach given the same budget and patience; they are not claimed here.

8.9 What batching costs in accuracy

The batched update is an approximation, so the honest question is how much fidelity it gives up. Figure 11 answers it on Noisy XOR — the benchmark of the original paper, deliberately chosen because its generating rule is known, so a failure to learn is unmistakable. Twenty clauses, $T = 15$, $s = 3.9$, 100 states, 2000 training examples with 40% label noise, three seeds per configuration, with the number of epochs raised as the batch grows so that every configuration sees the data a comparable number of times.

Batch sizes of roughly 8–64 keep the accuracy of the sequential algorithm while running one to two orders of magnitude faster on a GPU. At 200 the degradation is unambiguous. The practical recommendation the library documents follows directly: start at 16–64 for flat models and 8–32 for convolutional ones, and if accuracy lags, halve the batch or add epochs before touching T and s .

The cost of exactness is measurable too. At batch 32, `feedback_mode="sequential"` runs 7–28× slower than the batched update depending on model scale and device, because it performs 32 commits where the batched path performs one. That is the trade, stated in numbers: *batching buys one to two orders of magnitude and costs accuracy only once the batch grows past a few tens of examples.*




Figure 11. Noisy XOR accuracy against batch size (3 seeds). Batches up to about 50 land in the same band as the exact sequential algorithm; 200 degrades clearly. Seed-to-seed spread is wide on this task — which is itself the reason to read the individual points rather than only the means.

In short. Reformulating the update as event counting plus one binomial draw makes the cost of an update nearly independent of batch size — which is what turns a 114-second MNIST epoch into a 1.5-second one, and what makes a clause budget of thousands a routine choice rather than an overnight one.

9 Applications

9.1 What ships with the library

The repository includes eight runnable example scripts and five worked notebooks, chosen so that each one demonstrates a different part of the design rather than a different dataset.

The two segmentation notebooks are worth singling out. They reproduce the structure of published Tsetlin-machine segmentation work — convolutional regression [18] and a U-Net hybrid [19] — and the second builds a dense per-pixel classification head on top of `TsetlinMachine`. Neither required new model code: they are compositions of what the library already provides, which is the clearest evidence that the abstraction is at the right level.

9.2 Where a Tsetlin machine earns its place

Clinical decision support. This is the application domain the Tsetlin machine literature has pushed hardest, and for a reason: a clinician can audit a conjunction and a regulator can read it. Published work covers interpretable rule learning from medical text [22], premature ventricular contraction identification in ECG [21], organism prediction in peritoneal dialysis, bladder-cancer recurrence rule discovery and cervical-cancer classification with convolutional machines. What these share is a data regime — a few thousand examples, features that are already measurements with meaning — in which a thermometer-encoded Tsetlin machine is both competitive and legible, and a deep network is neither reliably better nor auditable.

`torchsetlin` contributes three things here: the encoders that turn clinical measurements into literals with names attached, so a rule comes out as `IF worst_perimeter >= 105.9 AND NOT mean_texture >= 21.8`; per-example explanations that are the model rather than a surrogate; and the ability to run the model inside the same PyTorch pipeline as the deep baseline it must be compared against.

Table 9. Worked applications shipped with torchsetlin. Timings are for the default arguments on an RTX 3090.

Example	Model	What it shows
<code>noisy_xor.py</code>	<code>TsetlinMachine</code>	the benchmark of the original paper; recovers the generating rule under 40% label noise
<code>tabular_breast_cancer.py</code>	<code>TsetlinMachine</code>	thermometer encoding of clinical measurements, rules printed in the original units
<code>regression_sine.py</code>	<code>RegressionTsetlinMachine</code>	a Tsetlin machine that emits a number
<code>mnist_flat.py</code>	<code>TsetlinMachine</code>	784 pixels as flat features, checkpoints, classification report
<code>shapes_conv.py</code>	<code>ConvTsetlinMachine</code>	1.000 test accuracy on every seed tried against 0.55–0.91 for a flat machine with the same clause budget — and what <code>position_encoding</code> costs on a translation-invariant task
<code>conv1d_ramps.py</code>	<code>Conv1dTsetlinMachine</code>	0.996 on noisy 1-D signals; the clauses print as the ramp they detect
<code>conv_regression_blobs.py</code>	<code>ConvRegressionTsetlinMachine</code>	reading a continuous quantity (a square’s side length) off an image: MAE 0.20, R^2 0.96
<code>mnist_conv.py</code>	<code>ConvTsetlinMachine</code> , <code>ConvCoalescedTsetlinMachine</code>	0.979 with 200 clauses per class; 0.985 with <code>--coalesced --clauses 1000 --T 500</code> , in less time
<i>notebooks</i>		
Iris	<code>TsetlinMachine</code>	a rule set small enough to read in full
MNIST, CIFAR-10	convolutional	clause patches, colour Booleanization, GPU throughput
Convolutional regression	<code>ConvRegressionTsetlinMachine</code>	working through the C-RTM segmentation literature ^[18]
CTM-UNet	<code>TsetlinMachine</code>	a dense per-pixel Tsetlin segmentation head, after ^[19]

Medical imaging. Convolutional and coalesced machines extend the same argument to images, and segmentation work ^[18,19] shows the family reaching per-pixel tasks. This is the setting that motivated torchsetlin: in a cardiac-MRI research group the surrounding stack is PyTorch, the volumes are large, and an interpretable model that cannot consume a `DataLoader` or run on the GPU next to the U-Net baseline is not a practical option. Making a convolutional Tsetlin machine an `nn.Module` is what makes that comparison a one-file experiment rather than a project.

Edge and hardware inference. Inference is an AND over included literals and a vote count, so a trained machine maps naturally onto digital logic. There is now a substantial literature on that mapping — asynchronous logic for edge inference ^[15], in-memory computing on Y-Flash cells ^[16], and 65 nm keyword-spotting accelerators ^[17]. torchsetlin is a *training* framework, not a deployment target, but the learned state it exports is exactly the artefact those pipelines consume: an integer matrix and a set of weights, in a `state_dict`, with a documented layout.

Signals and time series. `Conv1dTsetlinMachine` applies the same patch mechanism to sequences, where a clause becomes a motif detector that fires wherever the motif occurs. Combined with thermometer encoding of the signal amplitude, the learned clauses are legible as pictures of the waveform they detect — which is how the shipped `conv1d_ramps.py` example prints them.

Tabular and scientific data. The oldest and least glamorous case, and still the strongest: datasets with tens to hundreds of meaningful columns, where the question is often *which combinations of measurements matter* rather than what the accuracy is. Global feature importance aggregated back over thermometer bits answers that question in closed form, with no permutation testing and no surrogate.

In short. The applications that suit a Tsetlin machine are the ones where somebody has to read the model. `torchsetlin`'s contribution to them is prosaic and decisive: the model now lives in the same framework as everything else the project already uses.

10 Limitations and future development

10.1 Known limitations

Batched feedback is an approximation. The default path trades fidelity for throughput, and the trade is real: at batch 200 on Noisy XOR the degradation is unambiguous (Section 8.9). The exact algorithm is always one argument away, and the characterisation here is empirical — there is no proof that the batched update converges to the same fixed point as the sequential one, and establishing the conditions under which it does is open work.

Dense tensors, not bit-packed kernels. Literals are floats in a matmul, not bits in a machine word. A hand-written kernel can pack 64 literals per word and exploit clause sparsity through indexing [12]; `torchsetlin` does neither. The consequence is measurable and one-directional: on a *small* machine on a CPU, the reference C implementations are faster, and `torchsetlin`'s advantage only appears with model size and a GPU. The cached include mask also costs $2CF$ floats of memory that a packed representation would not need.

The commit is dominated by one operation. `torch.binomial` over the $(C, 2F)$ state is 43% of the CPU update and the single largest GPU kernel in the commit. It is a generic sampler being asked to do a very specific job; a fused kernel that samples and updates in one pass, or an approximation valid in the large-count regime, would move the whole library.

Variants not implemented. The Tsetlin machine literature has grown faster than any one package can follow. Absent from `torchsetlin` today: graph Tsetlin machines [14], TM composites [20], contracting machines with absorbing automata [24], the sparse machine with active literals [25], probabilistic and uncertainty-quantifying variants [26], and autoencoders. Each of these is expressible in the current architecture — most are an output layer and a feedback policy — but none is written.

No comparison against other implementations. Section 8 compares this library on a CPU against itself on a GPU. It does not establish how `torchsetlin` compares against `tmu` or the CUDA packages at matched configurations, on either speed or accuracy. That comparison is worth doing carefully, with matched Booleanization and matched hyper-parameters, and it has not been done here.

Interpretability coverage is uneven. Rule extraction, clause activity, and global importance work for every model. Local feature importance is weighted by literal values only for flat models;

for convolutional models the per-example attribution is computed over clause weights and inclusion without the per-patch literal gating, which makes it coarser than its flat counterpart.

Single-device, single-process. There is no multi-GPU or distributed training path today.

Early-stage API. Version 0.1.x is classified alpha. The core — `update`, `forward`, the buffer layout — is stable in intent, but the surface may still move, and checkpoints are tied to the state layout rather than a versioned format.

10.2 Where it goes next

A fused feedback kernel. The measurement points at exactly one target. A Triton or CUDA kernel that draws the binomial counts and applies the clamped update in a single pass over the state would cut the largest term in both the CPU and GPU profiles, and it fits the existing architecture as a drop-in replacement for one function in `functional.py`.

Data-parallel training, almost for free. This one falls out of the design. The feedback accumulator holds *additive* event counts: combining the counts from two chunks is addition, and nothing about that changes if the chunks were computed on different devices. Data-parallel training over n GPUs is therefore an all-reduce of four count tensors before a single commit — mathematically identical to a single-device update at n times the batch size, with no approximation introduced beyond the one batching already makes. The same property makes gradient-free federated training natural: what is exchanged is a count matrix, not a model.

A bit-packed inference path. Training benefits from dense float matmuls; inference does not need them. A packed representation for `predict` — and an export path to the integer artefacts that the hardware literature ^[15,16,17] consumes — would close the loop between this framework and deployment.

More of the literature. In rough order of how well each fits the current architecture: clause-size and sparsity-oriented variants (a feedback-policy change), TM composites (a wrapper over several trained machines), absorbing automata (a state-transition change in `apply_feedback`), and graph machines (a new input structure, in the same place convolution lives).

Validation against the reference implementations. A test suite that checks `torchsetlin`'s sequential path against the reference C implementation example by example, on fixed seeds, would turn “faithful to the published algorithm” from a claim into a measurement. The batched path could then be characterised as a controlled deviation from that baseline rather than from the literature's reported numbers.

In short. The honest summary: `torchsetlin` wins on integration and on large-model GPU throughput, loses to hand-written kernels on small CPU workloads, approximates the classical algorithm by default, and implements a subset of a fast-moving literature. Every one of those is a stated, measured trade rather than a surprise waiting for a user.

11 Availability

Version	0.1.3 (17 September 2026)
Source	https://github.com/vlbthambawita/torchsetlin
Package	https://pypi.org/project/torchsetlin/ — <code>pip install torchsetlin</code>
Documentation	https://vlbthambawita.github.io/torchsetlin/
Licence	MIT
Python	3.9, 3.10, 3.11, 3.12
Required	<code>torch >= 2.0</code> , <code>numpy >= 1.22</code> — no compiler, no CUDA toolkit, no C extension
Optional extras	<code>[viz]</code> matplotlib <code>[vision]</code> torchvision <code>[sklearn]</code> scikit-learn <code>[docs]</code> MkDocs Material <code>[dev]</code> the lot, plus pytest and ruff

11.1 Getting started

```
pip install torch          # pick the CPU or CUDA build from pytorch.org
pip install torchsetlin    # + optional extras: [viz] [vision] [sklearn] [all]
```

From source, with the development environment:

```
git clone https://github.com/vlbthambawita/torchsetlin && cd torchsetlin
pip install -e ".[dev]"
pytest -q                # the full suite
ruff check src tests      # exactly what CI runs
mkdocs serve              # the documentation site at localhost:8000
```

11.2 Reproducing the measurements

Every number in Section 8 comes from the driver in `benchmarks/`. It runs one suite at a time, merging its JSON after each, so a crash costs only the running suite, and then regenerates the figures, tables and the documentation page from those records.

```
./benchmarks/run_benchmarks.sh      # every suite, cpu + cuda:0 (~20-25 min)
./benchmarks/run_benchmarks.sh --quick  # ~1 min smoke test of the pipeline
./benchmarks/run_benchmarks.sh -s batch models # a subset
./benchmarks/run_benchmarks.sh --plots-only # re-render from results already on disk
```

The raw records land in `benchmarks/results/cpu_vs_gpu.json`. Replacing them with your own hardware's is the whole workflow — the derived page is regenerated, not hand-edited.

11.3 Quality gates

Continuous integration runs `ruff` and the full test suite on Python 3.9, 3.11 and 3.12, then deploys the documentation on pushes to `main`. Releases are cut from a `v*` tag through PyPI Trusted Publishing; the release workflow re-runs the suite and smoke-imports the built wheel in a clean virtual environment before uploading — a gate added after a `.gitignore` rule silently excluded a subpackage from the 0.1.0 wheel and made `import torchsetlin` fail for everyone who installed it.

11.4 Documentation

The documentation site is MkDocs Material with `mkdstrings` reading Google-style docstrings directly from the source, so the public API reference cannot drift from the code. It covers installation and a quick start; concept pages on the algorithm, Booleanization, batched versus sequential feedback, convolution and hyper-parameters; guides on training, evaluation, interpretability, visualisation, saving and extending; worked examples; and the generated CPU-versus-GPU benchmark page.

11.5 Citation

```
@software{torchsetlin,
  author = {Thambawita, Vajira},
  title = {torchsetlin: GPU-enabled, PyTorch-native Tsetlin machines},
  year = {2026},
  url = {https://github.com/vlbthambawita/torchsetlin}
}
```

Please also cite the Tsetlin machine papers whose algorithms you rely on — the reference list that follows is a reasonable starting point, and the concept pages in the documentation name the specific paper behind each option.

11.6 Contributing

Issues and pull requests are welcome at the repository. The most useful contributions, in order: a model variant from the literature implemented against the existing hooks (Section 5); a fused feedback kernel; validation against the reference implementations; and applications — a worked notebook on a real dataset is worth more to the next user than another synthetic benchmark.

References

- [1] O.-C. Granmo. *The Tsetlin Machine — A Game Theoretic Bandit Driven Approach to Optimal Pattern Recognition with Propositional Logic*. arXiv:1804.01508, 2018.
- [2] O.-C. Granmo, S. Glimsdal, L. Jiao, M. Goodwin, C. W. Omlin, G. T. Berge. *The Convolutional Tsetlin Machine*. arXiv:1905.09688, 2019.
- [3] A. Phoulady, O.-C. Granmo, S. R. Gorji, H. A. Phoulady. *The Weighted Tsetlin Machine: Compressed Representations with Weighted Clauses*. arXiv:1911.12607, 2019.
- [4] C. D. Blakely, O.-C. Granmo. *Closed-Form Expressions for Global and Local Interpretation of Tsetlin Machines*. arXiv:2007.13885, 2020.
- [5] S. Glimsdal, O.-C. Granmo. *Coalesced Multi-Output Tsetlin Machines with Clause Sharing*. arXiv:2108.07594, 2021.
- [6] K. D. Abeyrathna et al. *Building Concise Logical Patterns by Constraining Tsetlin Machine Clause Size*. arXiv:2301.08190, 2023.
- [7] K. D. Abeyrathna, O.-C. Granmo, X. Zhang, L. Jiao, M. Goodwin. *The Regression Tsetlin Machine: A Novel Approach to Interpretable Non-linear Regression*. Philosophical Transactions of the Royal Society A, 378(2164), 2020.
- [8] J. Sharma, R. Yadav, O.-C. Granmo, L. Jiao. *Drop Clause: Enhancing Performance, Robustness and Pattern Recognition Capabilities of the Tsetlin Machine*. AAAI 2023; arXiv:2105.14506.
- [9] O.-C. Granmo. *An Introduction to Tsetlin Machines*. <https://tsetlinmachine.org>.

- [10] M. L. Tsetlin. *On the Behaviour of Finite Automata in Random Media*. Automation and Remote Control, 22:1210–1219, 1961.
- [11] K. D. Abeyrathna et al. *Massively Parallel and Asynchronous Tsetlin Machine Architecture*. ICML 2021; arXiv:2009.04861.
- [12] S. R. Gorji et al. *Increasing the Inference and Learning Speed of Tsetlin Machines with Clause Indexing*. arXiv:2004.03188, 2020.
- [13] Halenka et al. *Exploring the Effects of Hyperdimensional Vectors for Tsetlin Machines*. arXiv:2406.02648, 2024.
- [14] O.-C. Granmo et al. *The Tsetlin Machine Goes Deep: Logical Learning and Reasoning with Graphs*. arXiv:2507.14874, 2025.
- [15] A. Wheeldon, R. Shafik, A. Yakovlev et al. *Low-Latency Asynchronous Logic Design for Inference at the Edge*. arXiv:2012.03402, 2020.
- [16] *IMPACT: In-Memory Computing Architecture Based on Y-Flash Technology for Coalesced Tsetlin Machine Inference*. arXiv:2412.05327, 2024.
- [17] *TsetlinKWS: 65 nm State-Driven Convolutional Tsetlin Machine Accelerator for Keyword Spotting*. arXiv:2510.24282, 2025.
- [18] K. D. Abeyrathna, O.-C. Granmo, M. Goodwin. *Convolutional Regression Tsetlin Machine: An Interpretable Approach to Convolutional Regression*. ICMLT 2021. doi:10.1145/3468891.3468901
- [19] *CTM-UNet: Integrating Convolutional Tsetlin Machines with U-Net Architecture for Semantic Segmentation*. ISTM 2025.
- [20] O.-C. Granmo. *TMComposites: Plug-and-Play Collaboration Between Specialized Tsetlin Machines*. arXiv:2309.04801, 2023.
- [21] *Interpretable Tsetlin Machine-based Premature Ventricular Contraction Identification from ECG*. arXiv:2301.10181, 2023.
- [22] G. T. Berge, O.-C. Granmo, T. Tveit, M. Goodwin, L. Jiao, B. Matheussen. *Using the Tsetlin Machine to Learn Human-Interpretable Rules for High-Accuracy Text Categorization with Medical Applications*. arXiv:1809.04547, 2018.
- [23] A. Paszke et al. *PyTorch: An Imperative Style, High-Performance Deep Learning Library*. NeurIPS 2019.
- [24] *Contracting Tsetlin Machine with Absorbing Automata*. arXiv:2310.11481, 2023.
- [25] *The Sparse Tsetlin Machine: Sparse Representation with Active Literals*. arXiv:2405.02375, 2024.
- [26] *The Probabilistic Tsetlin Machine: A Novel Approach to Uncertainty Quantification*. arXiv:2410.17851, 2024.

arXiv identifiers and DOIs are given so that every entry can be resolved directly. Where an author list is abbreviated or omitted, the title and identifier are authoritative. A BibTeX version of this list ships alongside the source of this document as `references.bib`.